

Threads

Mārtiņš Leitass

Task Scheduling

Modern operating systems use preemptive multitasking to allocate CPU time to applications. There are two types of tasks that can be scheduled for execution:

- Processes: A process is an area of memory that contains both code and data. A process has a thread of execution that is scheduled to receive CPU time slices.
- Thread: A thread is a scheduled execution of a process. Concurrent threads are possible. All threads for a process share the same data memory but may be following different paths through a code section.

Concurrent programming

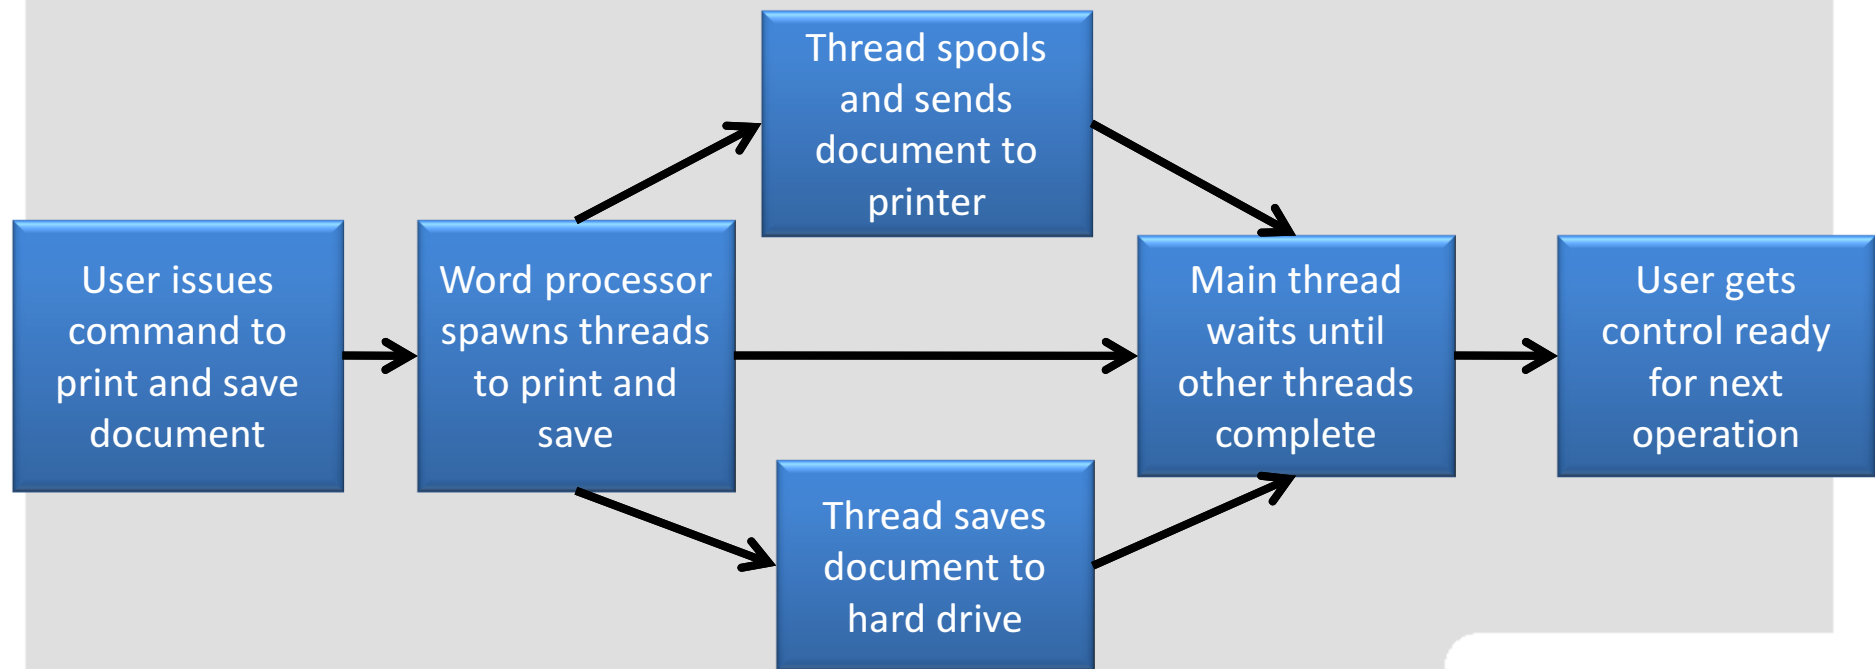
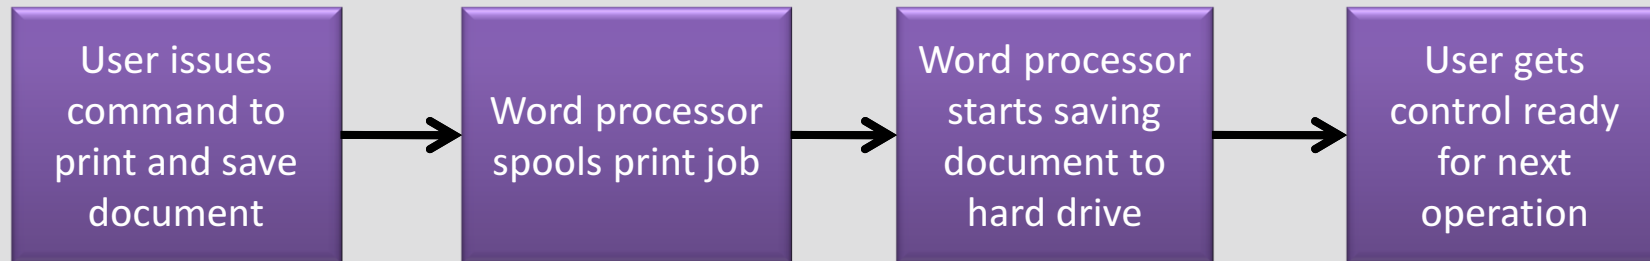
- with a single-processor CPU, only one thread is executing at any given time
- on multi-processor systems several threads are actually executing at the same time (physical concurrency)
- **multi-threading** occurs when concurrency exists among threads running in a single process (also referred to as multi-tasking)

Why Threading Matters

To execute a program as quickly as possible, you must avoid performance bottlenecks. Some of these bottlenecks are:

- Resource Contention: Two or more tasks waiting for exclusive use of a resource
- Blocking I/O operations: Doing nothing while waiting for disk or network data transfers
- Underutilization of CPUs: A single-threaded application uses only a single CPU

Concurrent programming



Java Threads

- Java provides support for multi-threading as a part of the language
- support centers on the:
 - `java.lang.Thread` class
 - `java.lang.Runnable` interface
 - `java.lang.Object` methods `wait()`, `notify()`, and `notifyAll`
 - `synchronized` keyword
- every Java program has at least one thread which is executed when `main()` is invoked
- all user-level threads are explicitly constructed and started from the main thread or by a thread originally started from `main()`
- when the last **user** thread completes any **daemon** threads are stopped and the application stops
- a thread's default daemon status is the same as that of thread creating it
 - you can check the daemon status using **`isDaemon()`**
 - you can set the daemon status using **`setDaemon()`**.
- You cannot change a thread's status after it has been started


Creating the Thread

```
1  public class ThreadTester {  
2      public static void main(String args[]) {  
3          HelloRunner r = new HelloRunner();  
4          Thread t = new Thread(r);  
5          t.start();  
6      }  
7  }  
8  class HelloRunner implements Runnable {  
9      int i;  
10     public void run() {  
11         i = 0;  
12         while (true) {  
13             System.out.println("Hello " + i++);  
14             if ( i == 50 ) {  
15                 break;  
16             }  
17         }  
18     }  
19 }
```

Creating the Thread

- Multithreaded programming has these characteristics:
 - Multiple threads are from one Runnable instance.
 - Threads share the same data and code.
- For example:

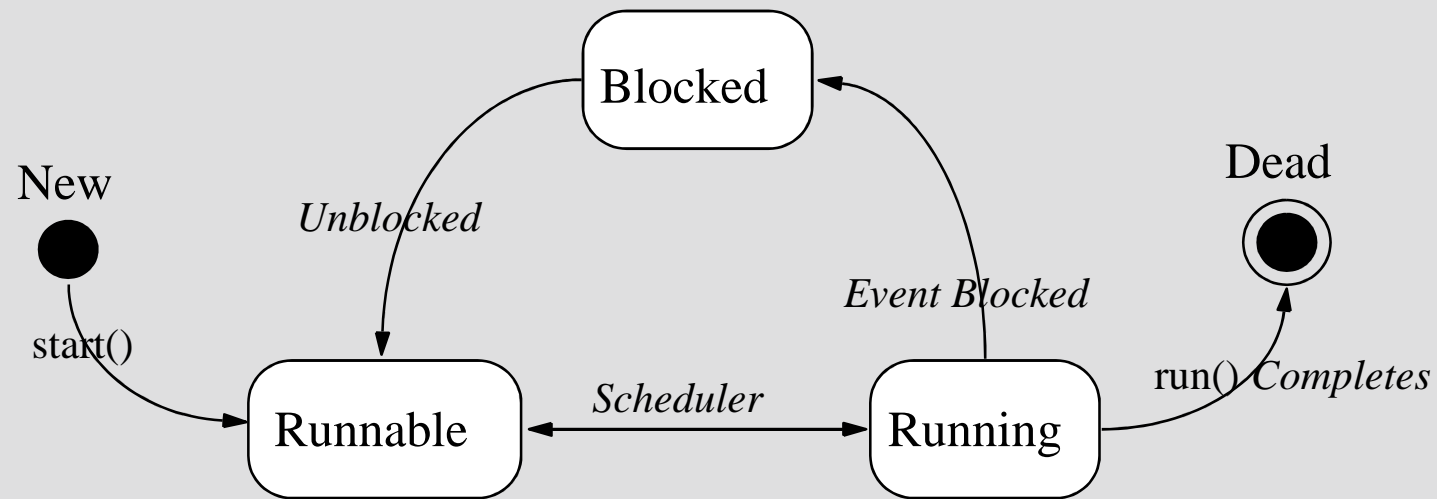
```
Thread t1 = new Thread(r);  
Thread t2 = new Thread(r);
```



Starting the Thread

- Use the `start` method.
- Place the thread in a runnable state.

Thread Scheduling



Thread Scheduling Example

```
1  public class Runner implements Runnable {  
2      public void run() {  
3          while (true) {  
4              // do lots of interesting stuff  
5              // ...  
6              // Give other threads a chance  
7              try {  
8                  Thread.sleep(10);  
9              } catch (InterruptedException e) {  
10                 // This thread's sleep was interrupted  
11                 // by another thread  
12             }  
13         }  
14     }  
15 }
```

Terminating a Thread

```
1  public class Runner implements Runnable {  
2      private boolean timeToQuit=false;  
3  
4      public void run() {  
5          while ( ! timeToQuit ) {  
6              // continue doing work  
7          }  
8          // clean up before run() ends  
9      }  
10  
11     public void stopRunning() {  
12         timeToQuit=true;  
13     }  
14 }
```

Terminating a Thread

```
1  public class ThreadController {  
2      private Runner r = new Runner();  
3      private Thread t = new Thread(r);  
4  
5      public void startThread() {  
6          t.start();  
7      }  
8  
9      public void stopThread() {  
10         // use specific instance of Runner  
11         r.stopRunning();  
12     }  
13 }
```

Basic Control of Threads

- Test threads:

`isAlive()`

- Access thread priority:

`getPriority()`

`setPriority()`

- Put threads on hold:

`Thread.sleep()` *// static method*

`join()`

`Thread.yield()` *// static method*

The join Method

```
1  public static void main(String[] args) {
2      Thread t = new Thread(new Runner());
3      t.start();
4      ...
5      // Do stuff in parallel with the other thread for a while
6      ...
7      // Wait here for the other thread to finish
8      try {
9          t.join();
10     } catch (InterruptedException e) {
11         // the other thread came back early
12     }
13     ...
14     // Now continue in this thread
15     ...
16 }
```

Other Ways to Create Threads

```
1  public class MyThread extends Thread {  
2      public void run() {  
3          while ( true ) {  
4              // do lots of interesting stuff  
5              try {  
6                  Thread.sleep(100);  
7              } catch (InterruptedException e) {  
8                  // sleep interrupted  
9              }  
10         }  
11     }  
12  
13     public static void main(String args[]) {  
14         Thread t = new MyThread();  
15         t.start();  
16     }  
17 }
```


Selecting a Way to Create Threads

- Implement Runnable:
 - Better object-oriented design
 - Single inheritance
 - Consistency
- Extend Thread:
Simpler code

Out-of-Order Execution


- Operations performed in one thread may not appear to execute in order if you observe the results from another thread.
 - Code optimization may result in out-of-order operation.
 - Threads operate on cached copies of shared variables.
- To ensure consistent behavior in your threads, you must synchronize their actions.
 - You need a way to state that an action happens before another.
 - You need a way to flush changes to shared variables back to main memory.

Using the synchronized Keyword


```
1  public class MyStack {  
2  
3      int idx = 0;  
4      char [] data = new char[6];  
5  
6      public void push(char c) {  
7          data[idx] = c;  
8          idx++;  
9      }  
10  
11     public char pop() {  
12         idx--;  
13         return data[idx];  
14     }  
15 }
```

The Object Lock Flag

- Every object has a flag that is a type of *lock flag*.
- The `synchronized` enables interaction with the lock flag.

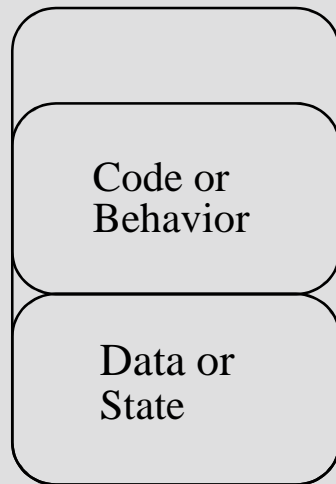


The Object Lock Flag



The Object Lock Flag

Object this
lock flag missing



Another thread, trying to
execute synchronized(this)

Waiting for
object lock

```
public char pop() {  
    synchronized (this) {  
        idx--;  
        return data[idx];  
    }  
}
```

Releasing the Lock Flag

The lock flag is released in the following events:

- Released when the thread passes the end of the synchronized code block
- Released automatically when a break, return, or exception is thrown by the synchronized code block

Object Monitor Locking

- `synchronized` methods use the monitor for the this object.
- `static synchronized` methods use the classes' monitor.
- `synchronized` blocks must specify which object's monitor to lock or unlock.
- `synchronized` blocks can be nested.

Using synchronized– Putting It Together

- *All* access to delicate data should be synchronized.
- Delicate data protected by synchronized should be private.

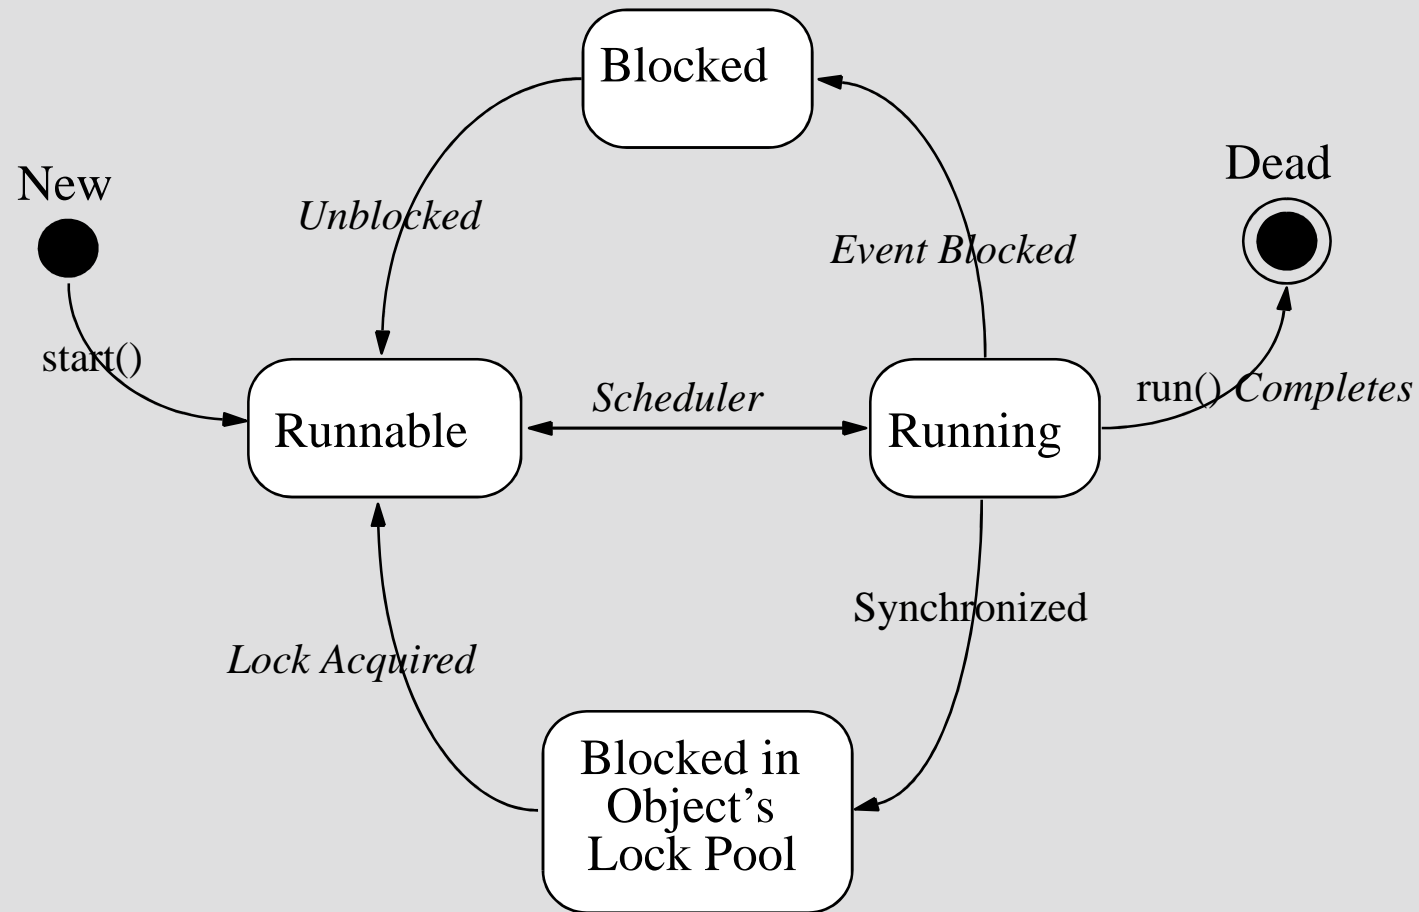
Using synchronized– Putting It Together

The following two code segments are equivalent:

```
public void push(char c) {  
    synchronized(this) {  
        // The push method code  
    }  
}
```

```
public synchronized void push(char c) {  
    // The push method code  
}
```

Thread State Diagram With Synchronization



The volatile Keyword

A field may have the volatile modifier applied to it:

- Reading or writing a volatile field will cause a thread to synchronize its working memory with main memory.
- volatile does not mean atomic.
 - If i is volatile, i++ is still not a thread-safe operation.

```
public volatile int i;
```

Deadlock

A deadlock has the following characteristics:

- It is two threads, each waiting for a lock from the other.
- It is not detected or avoided.
- Deadlock can be avoided by:
 - Deciding on the order to obtain locks
 - Adhering to this order throughout
 - Releasing locks in reverse order

Thread Interaction – wait and notify

- Scenario:
Consider yourself and a cab driver as two threads.
- The problem:
How do you determine when you are at your destination?
- The solution:
 - You notify the cab driver of your destination and relax.
 - The driver drives and notifies you upon arrival at your destination.

Thread Interaction

Thread interactions include:

- The wait and notify methods
- The pools:
 - Wait pool
 - Lock pool

Monitor Model for Synchronization

- Leave shared data in a consistent state.
- Ensure programs cannot deadlock.
- Do not put threads expecting different notifications in the same wait pool.

The Producer Class

```
1  package mod13;
2
3  public class Producer implements Runnable {
4      private SyncStack theStack;
5      private int num;
6      private static int counter = 1;
7
8      public Producer (SyncStack s) {
9          theStack = s;
10         num = counter++;
11     }
12
```

The Producer Class

```
13 public void run() {
14     char c;
15
16     for (int i = 0; i < 200; i++) {
17         c = (char)(Math.random() * 26 + 'A');
18         theStack.push(c);
19         System.out.println("Producer" + num + ": " + c);
20         try {
21             Thread.sleep((int)(Math.random() * 300));
22         } catch (InterruptedException e) {
23             // ignore it
24         }
25     }
26 } // END run method
27
28 } // END Producer class
```

The Consumer Class

```
1  package mod13;
2
3  public class Consumer implements Runnable {
4      private SyncStack theStack;
5      private int num;
6      private static int counter = 1;
7
8      public Consumer (SyncStack s) {
9          theStack = s;
10         num = counter++;
11     }
12
```

The Consumer Class

```
13 public void run() {  
14     char c;  
15     for (int i = 0; i < 200; i++) {  
16         c = theStack.pop();  
17         System.out.println("Consumer" + num + ": " + c);  
18  
19         try {  
20             Thread.sleep((int)(Math.random() * 300));  
21         } catch (InterruptedException e) {  
22             // ignore it  
23         }  
24     }  
25 } // END run method  
26
```

The SyncStackClass

This is a sketch of the SyncStack class:

```
public class SyncStack {  
  
    private List<Character> buffer = new ArrayList<Character>(400);  
  
    public synchronized char pop() {  
        // pop code here  
    }  
  
    public synchronized void push(char c) {  
        // push code here  
    }  
}
```

The popMethod

```
9  public synchronized char pop() {  
10      char c;  
11      while (buffer.size() == 0) {  
12          try {  
13              this.wait();  
14          } catch (InterruptedException e) {  
15              // ignore it...  
16          }  
17      }  
18      c = buffer.remove(buffer.size()-1);  
19      return c;  
20  }  
21
```

The pushMethod

```
22 public synchronized void push(char c) {  
23     this.notify();  
24     buffer.add(c);  
25 }
```

The SyncTestClass

```
1  package mod13;
2  public class SyncTest {
3      public static void main(String[] args) {
4          SyncStack stack = new SyncStack();
5          Producer p1 = new Producer(stack);
6          Thread prodT1 = new Thread (p1);
7          prodT1.start();
8          Producer p2 = new Producer(stack);
9          Thread prodT2 = new Thread (p2);
10         prodT2.start();
11
12         Consumer c1 = new Consumer(stack);
13         Thread consT1 = new Thread (c1);
14         consT1.start();
15         Consumer c2 = new Consumer(stack);
16         Thread consT2 = new Thread (c2);
17         consT2.start();
18     }
19 }
```


The SyncTestClass

Producer2: F
Consumer1: F
Producer2: K
Consumer2: K
Producer2: T
Producer1: N
Producer1: V
Consumer2: V
Consumer1: N
Producer2: V
Producer2: U
Consumer2: U
Consumer2: V
Producer1: F
Consumer1: F
Producer2: M
Consumer2: M
Consumer2: T

Methods to Avoid

Some Thread methods should be avoided:

- `setPriority(int)` and `getPriority()`
 - Might not have any impact or may cause problems
- The following methods are deprecated and should never be used:
 - `destroy()`
 - `resume()`
 - `suspend()`
 - `stop()`