

# JDBC and Database Programming in Java

# Agenda


- Overview of JDBC
- JDBC APIs

# JDBC Overview


# JDBC Goals

- SQL-Level
- 100% Pure Java
- Keep it simple
- High-performance
- Leverage existing database technology
  - why reinvent the wheel?
- Use strong, static typing wherever possible
- Use multiple methods to express multiple functionality

# JDBC Ancestry



# JDBC Architecture

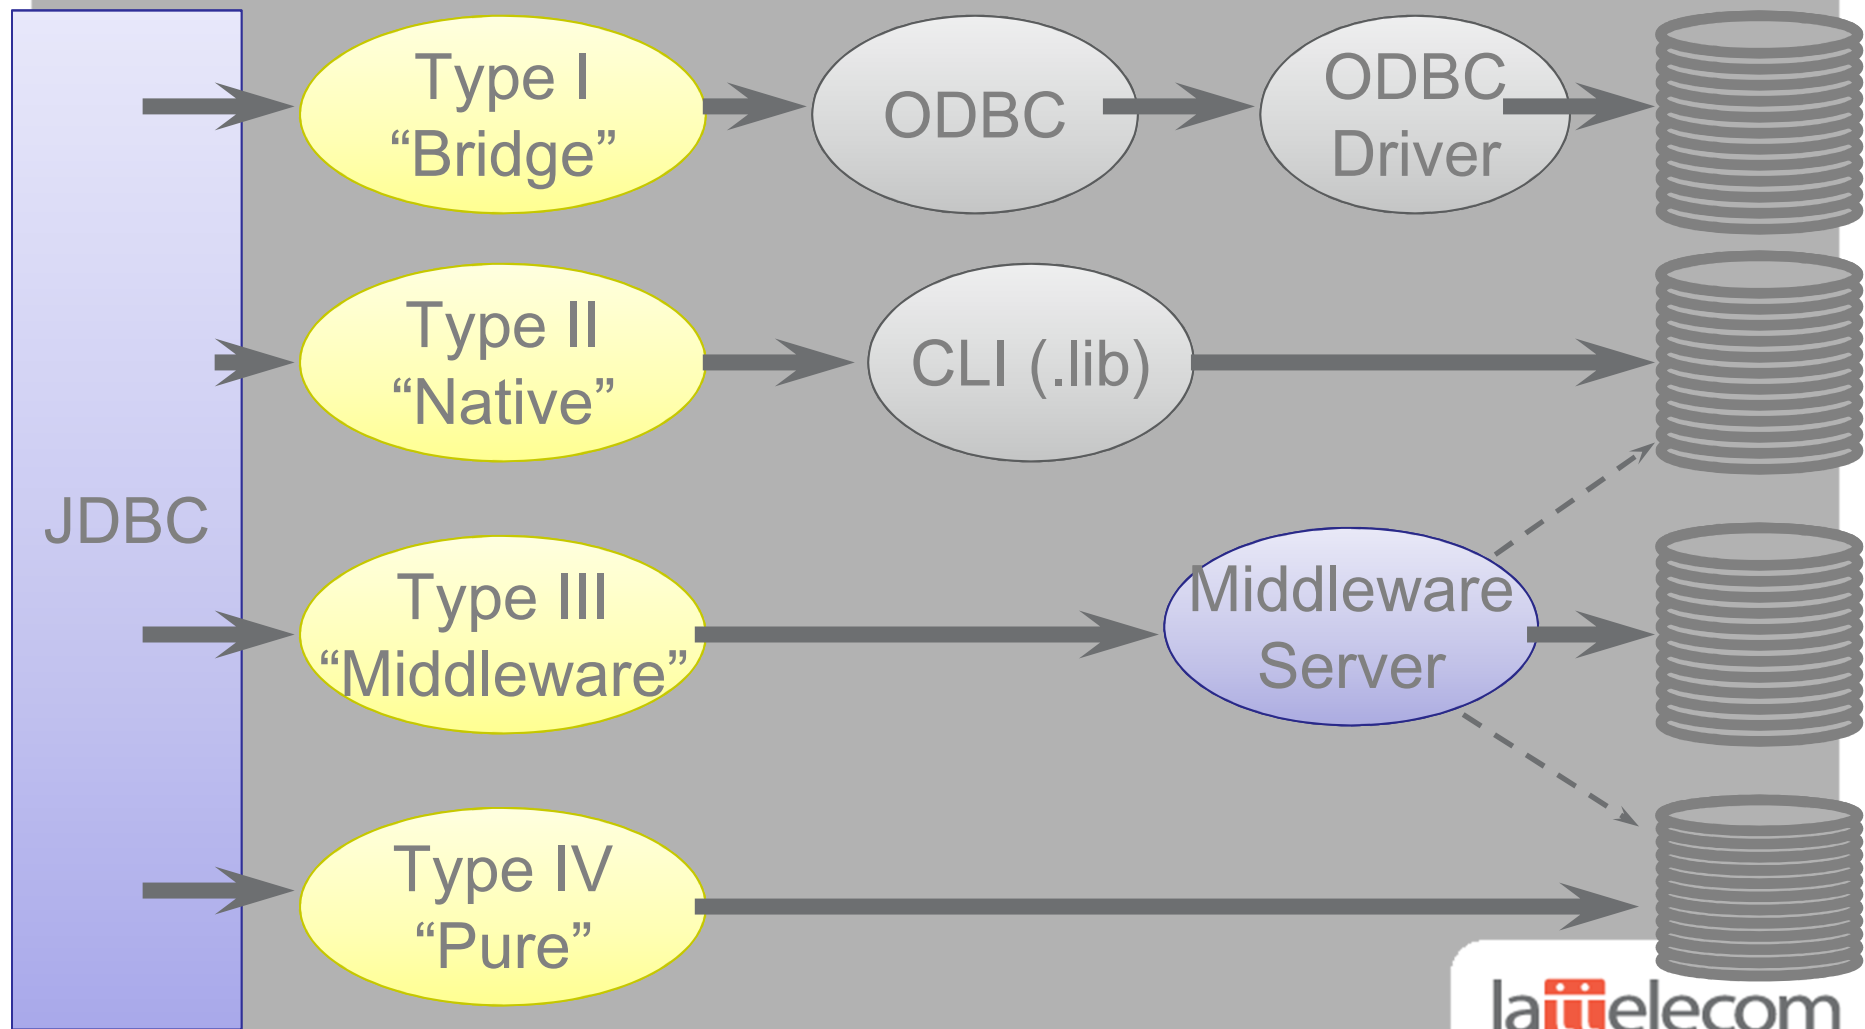


- Java code calls JDBC library
- JDBC loads a *driver*
- Driver talks to a particular database
- Can have more than one driver -> more than one database
- Ideal: can change database engines without changing any application code

# JDBC Drivers

- Type I: “Bridge”
- Type II: “Native”
- Type III: “Middleware”
- Type IV: “Pure”

# JDBC Drivers (Fig.)





# Type I Drivers

- Use bridging technology
- Requires installation/configuration on client machines
- Not good for Web
- e.g. ODBC Bridge

# Type II Drivers

- Native API drivers
- Requires installation/configuration on client machines
- Used to leverage existing CLI libraries
- Usually not thread-safe
- Mostly obsolete now

# Type III Drivers

- Calls middleware server, usually on database host
- Very flexible -- allows access to multiple databases using one driver
- Only need to download one driver
- But it's another server application to install and maintain

# Type IV Drivers

- 100% Pure Java -- the Holy Grail
- Use Java networking libraries to talk directly to database engines
- Only disadvantage: need to download a new driver for each database engine
- e.g. Oracle

# Related Technologies

- ODBC
  - Requires configuration (odbc.ini)
- RDO, ADO
  - Requires Win32
- OODB
  - e.g. ObjectStore from ODI
- JPA, ORM
  - maps objects to tables transparently (more or less)

# JDBC APIs

# java.sql

- JDBC is implemented via classes in the java.sql package

# Loading a Driver Directly

```
Driver d = new foo.bar.MyDriver();
```

```
Connection c = d.connect(...);
```

- Not recommended, use DriverManager instead
- Useful if you know you want a particular driver



# DriverManager

- DriverManager tries all the drivers
- Uses the first one that works
- When a driver class is first loaded, it registers itself with the DriverManager
- Therefore, to register a driver, just load it!

# Registering a Driver

- statically load driver

```
Class.forName( "foo.bar.MyDriver" );
```

```
Connection c =
```

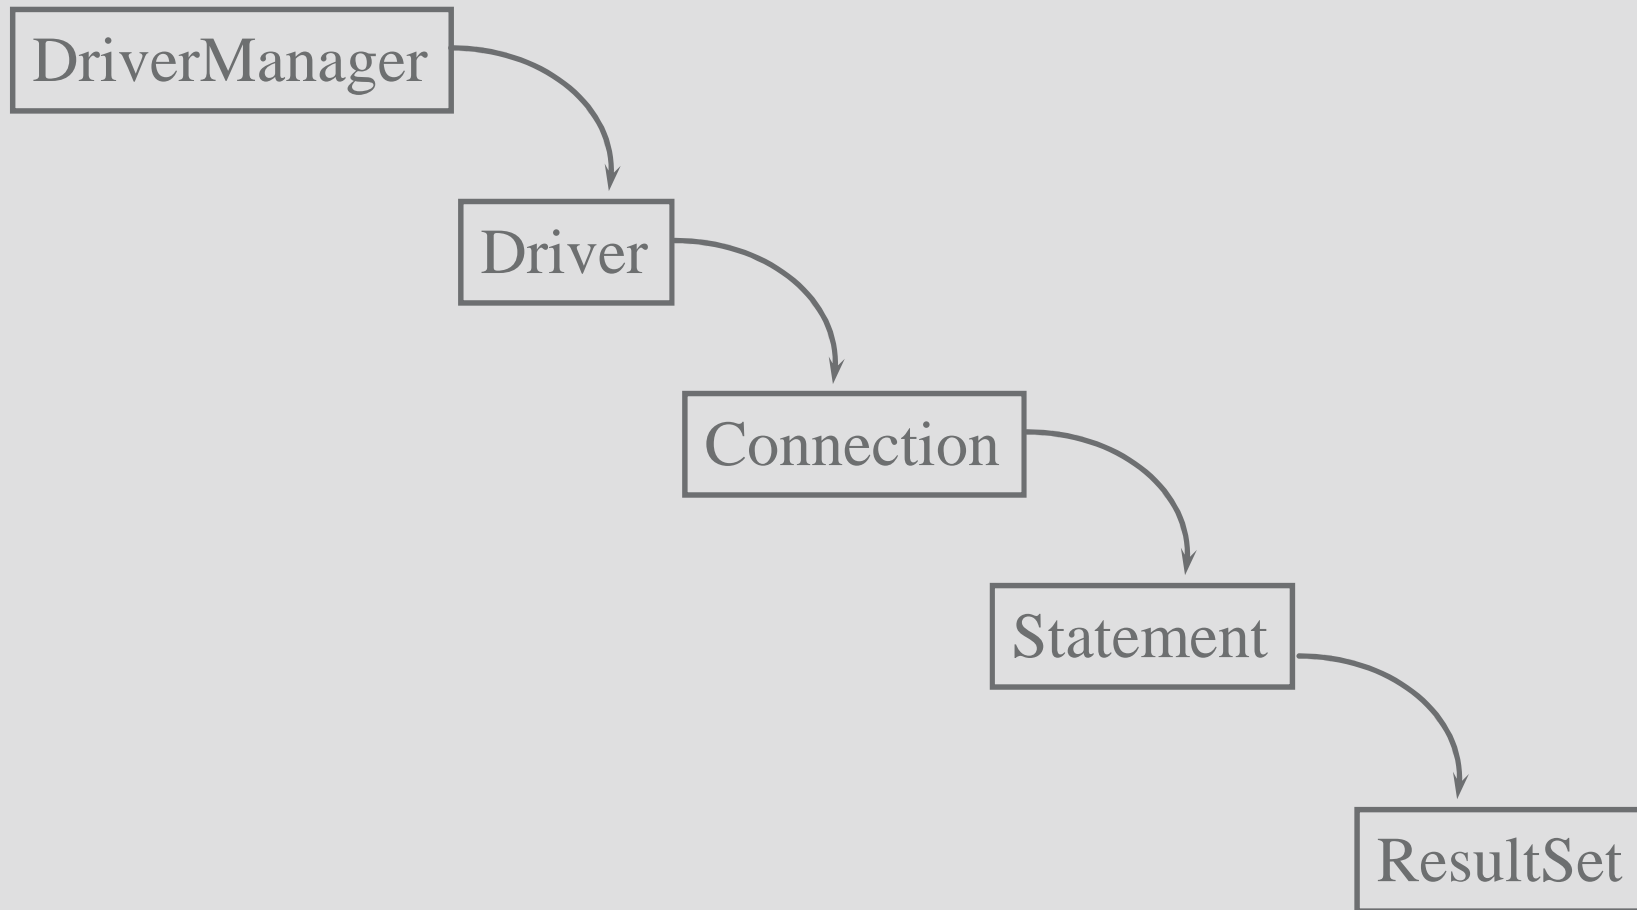
```
    DriverManager.getConnection( ... );
```

- or use the `jdbc.drivers` system property

# JDBC Object Classes

- DriverManager
  - Loads, chooses drivers
- Driver
  - connects to actual database
- Connection
  - a series of SQL statements to and from the DB
- Statement
  - a single SQL statement
- ResultSet
  - the records returned from a Statement

# JDBC Class Usage



# JDBC URLs

`jdbc:subprotocol:source`

- each driver has its own subprotocol
- each subprotocol has its own syntax for the source

`jdbc:odbc:DataSource`

- e.g. `jdbc:odbc:Northwind`

`jdbc:mysql://host[:port]/database`

- e.g.

`jdbc:mysql://foo.nowhere.com:4333/accounting`

# DriverManager

`Connection getConnection`

`(String url, String user, String password)`

- Connects to given JDBC URL with given user name and password
- Throws `java.sql.SQLException`
- returns a `Connection` object

# Connection

- A Connection represents a session with a specific database.
- Within the context of a Connection, SQL statements are executed and results are returned.
- Can have multiple connections to a database
  - NB: Some drivers don't support serialized connections
  - Fortunately, most do (now)
- Also provides “metadata” -- information about the database, tables, and fields
- Also methods to deal with transactions

# Obtaining a Connection

```
String url    = "jdbc:odbc:Northwind";  
try {  
    Class.forName ("sun.jdbc.odbc.JdbcOdbcDriver");  
    Connection con = DriverManager.getConnection(url);  
}  
catch (ClassNotFoundException e)  
    { e.printStackTrace(); }  
catch (SQLException e)  
    { e.printStackTrace(); }
```



# Connection Methods

**Statement createStatement()**

- returns a new Statement object

**PreparedStatement prepareStatement(String sql)**

- returns a new PreparedStatement object

**CallableStatement prepareCall(String sql)**

- returns a new CallableStatement object

- Why all these different kinds of statements?  
Optimization.

# Statement

- A Statement object is used for executing a static SQL statement and obtaining the results produced by it.

# Statement Methods

`ResultSet executeQuery(String)`

- Execute a SQL statement that returns a single `ResultSet`.

`int executeUpdate(String)`

- Execute a SQL INSERT, UPDATE or DELETE statement. Returns the number of rows changed.

`boolean execute(String)`

- Execute a SQL statement that may return multiple results.

# ResultSet

- A ResultSet provides access to a table of data generated by executing a Statement.
- Only one ResultSet per Statement can be open at once.
- The table rows are retrieved in sequence.
- A ResultSet maintains a cursor pointing to its current row of data.
- The 'next' method moves the cursor to the next row.

# ResultSet Methods

- `boolean next()`
  - activates the next row
  - the first call to `next()` activates the first row
  - returns false if there are no more rows
- `void close()`
  - disposes of the `ResultSet`
  - allows you to re-use the `Statement` that created it
  - automatically called by most `Statement` methods

# ResultSet Methods

- *Type* get *Type*(int columnIndex)
  - returns the given field as the given type
  - fields indexed starting at 1 (not 0)
- *Type* get *Type*(String columnName)
  - same, but uses name of field
  - less efficient
- int findColumn(String columnName)
  - looks up column index given column name

# ResultSet Methods

- String getString(int columnIndex)
- boolean getBoolean(int columnIndex)
- byte getByte(int columnIndex)
- short getShort(int columnIndex)
- int getInt(int columnIndex)
- long getLong(int columnIndex)
- float getFloat(int columnIndex)
- double getDouble(int columnIndex)
- Date getDate(int columnIndex)
- Time getTime(int columnIndex)
- Timestamp getTimestamp(int columnIndex)

# ResultSet Methods

- String getString(String columnName)
- boolean getBoolean(String columnName)
- byte getByte(String columnName)
- short getShort(String columnName)
- int getInt(String columnName)
- long getLong(String columnName)
- float getFloat(String columnName)
- double getDouble(String columnName)
- Date getDate(String columnName)
- Time getTime(String columnName)
- Timestamp getTimestamp(String columnName)



# Sample Database

Employee ID	Last Name	First Name
1	Davolio	Nancy
2	Fuller	Andrew
3	Leverling	Janet
4	Peacock	Margaret
5	Buchanan	Steven

# SELECT Example

```
Connection con =  
    DriverManager.getConnection(url,  
        "alex", "8675309");  
  
Statement st = con.createStatement();  
  
ResultSet results =  
    st.executeQuery("SELECT EmployeeID,  
        LastName, FirstName FROM Employees");
```

## SELECT Example (Cont.)

```
while (results.next()) {  
    int id = results.getInt(1);  
    String last = results.getString(2);  
    String first = results.getString(3);  
    System.out.println(" " + id + ": " +  
        first + " " + last);  
}  
st.close();  
con.close();
```

# Mapping Java Types to SQL Types

<u>SQL type</u>	<u>Java Type</u>
CHAR, <u>VARCHAR</u> , LONGVARCHAR	String
<u>NUMERIC</u> , DECIMAL	java.math.BigDecimal
BIT	boolean
TINYINT	byte
SMALLINT	short
INTEGER	int
BIGINT	long
REAL	float
FLOAT, <u>DOUBLE</u>	double
BINARY, <u>VARBINARY</u> , LONGVARBINARY	byte[]
DATE	java.sql.Date
TIME	java.sql.Time
TIMESTAMP	java.sql.Timestamp

# Modifying the Database

- use `executeUpdate` if the SQL contains “INSERT” or “UPDATE”
- Why isn't it smart enough to parse the SQL?  
Optimization.
- `executeUpdate` returns the number of rows modified
- `executeUpdate` also used for “CREATE TABLE” etc.  
(DDL)

# setAutoCommit

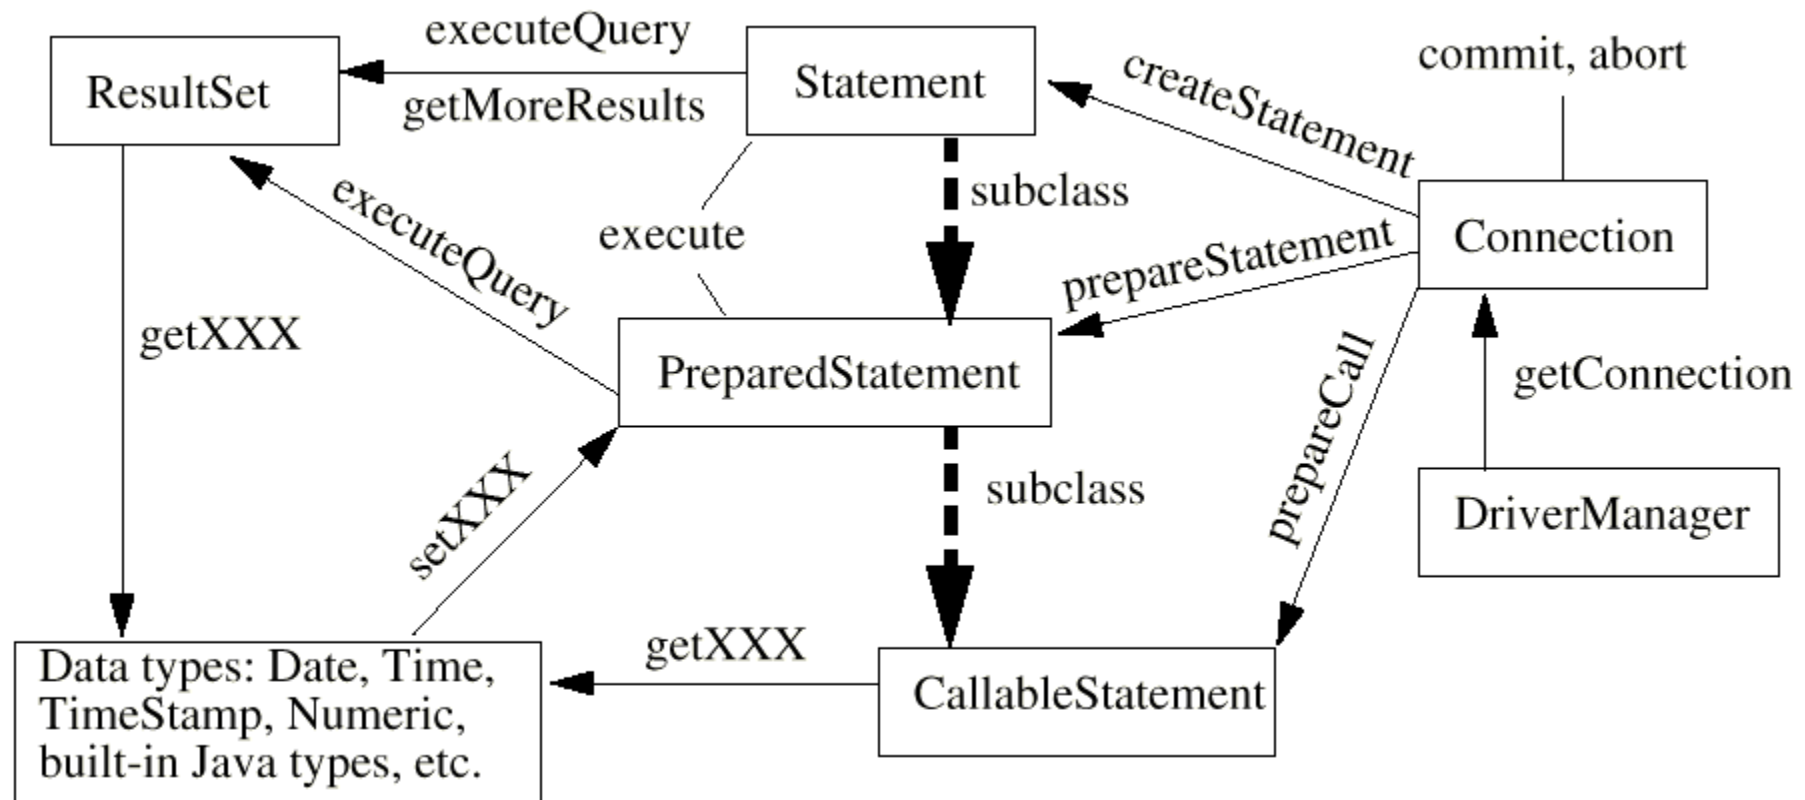
`Connection.setAutoCommit(boolean)`

- if *AutoCommit* is false, then every statement is added to an ongoing transaction
- you must explicitly commit or rollback the transaction using `Connection.commit()` and `Connection.rollback()`

# Optimized Statements

- Prepared Statements
  - SQL calls you make again and again
  - allows driver to optimize (compile) queries
  - created with `Connection.prepareStatement()`
- Stored Procedures
  - written in DB-specific language
  - stored inside database
  - accessed with `Connection.prepareCall()`

# JDBC Class Diagram



Whoa!



# Metadata

- Connection:
  - DatabaseMetaData getMetaData()
- ResultSet:
  - ResultSetMetaData getMetaData()

# ResultSetMetaData

- What's the number of columns in the ResultSet?
- What's a column's name?
- What's a column's SQL type?
- What's the column's normal max width in chars?
- What's the suggested column title for use in printouts and displays?
- What's a column's number of decimal digits?
- Does a column's case matter?
- Is the column a cash value?
- Will a write on the column definitely succeed?
- Can you put a NULL in this column?
- Is a column definitely not writable?
- Can the column be used in a where clause?
- Is the column a signed number?
- Is it possible for a write on the column to succeed?
- and so on...

# DatabaseMetaData

- What tables are available?
- What's our user name as known to the database?
- Is the database in read-only mode?
- If table correlation names are supported, are they restricted to be different from the names of the tables?
- and so on...

# JDBC 2.0

- Scrollable result set
- Batch updates
- Advanced data types
  - Blobs, objects, structured types
- Rowsets
  - Persistent JavaBeans
- JNDI
- Connection Pooling
- Distributed transactions via JTS

# Where to get more information

- Other training sessions
- Reese, *Database Programming with JDBC and Java* (O'Reilly)
- <http://java.sun.com/products/jdbc/>
- <http://docs.oracle.com/javase/tutorial/jdbc/index.html>